

Panel-Driven Research Quality Impact: Comparing Traditional and AI-Simulated Expert Review

Gio Della-Libera, VP Engineering
Business & Industry Copilots Agents 365

February 2026

Abstract

Does structured peer review simulation improve AI-assisted research quality, or merely add process overhead? We present a case study comparing research outputs from the same computational redistricting project under two conditions: traditional Claude-assisted writing (user-directed, no formal review) and panel-driven methodology (AI-simulated expert review with iterative refinement). Through within-project analysis of papers from the apportionment project, we observe that panel-driven research demonstrates substantially higher rigor across 10 quality dimensions: hypothesis clarity (+150%), experimental rigor (+150%), systematic testing (+400%), theory building (+150%), and innovation level (+150%), yielding an overall quality difference of +127%. While statistical inference is not possible given the small sample (N=1 panel-driven paper), we identify four mechanisms that may explain observed differences: systematic questioning, embracing negative results, standards elevation, and iteration forcing. Process tracing reveals that the panel acts as a scientific driver through iterative feedback, with the user's role shifting from director to facilitator. Breakthrough innovations (42% demographic threshold, edge-weighting methodology) emerged through panel-driven systematic exploration of failures. This case study generates hypotheses about how panel-driven research may elevate AI-assisted work, with implications for autonomous AI research agents and human-AI research collaboration. We propose a rigorous experimental design for future validation.

1 Introduction

AI-assisted research is proliferating rapidly—large language models now help draft papers, generate code, design experiments, and analyze results [36]. Yet a fundamental question remains under-explored: **Does structured peer review simulation improve AI-assisted research quality, or merely add process overhead?**

Two paradigms dominate current practice:

Traditional Claude-assisted research: The user directs the investigation, Claude executes instructions, and quality control relies on user judgment. This model treats the AI as a highly capable assistant responding to user direction.

Panel-driven research: Papers undergo AI-simulated expert review through an 8-stage life-cycle [9]—from draft through reviewer panel assembly, multi-round synthesis, systematic revision, and iterative recheck before submission. The panel drives the scientific investigation; the user facilitates execution.

The distinction is not merely procedural. In traditional workflows, users must possess sufficient domain expertise to direct the research, identify gaps, and demand systematic testing. Panel-

driven workflows shift this burden to the review process itself, raising a provocative hypothesis: **The panel may elevate research quality beyond what user direction alone achieves.**

We explore this question through a case study within a single computational redistricting project. The apportionment repository contains both traditional Claude-assisted papers (written in early 2026 with user direction) and panel-driven research (VRA compliance investigation conducted under systematic review). This within-project comparison controls for domain, author, AI model, and timeframe, isolating the effect of review methodology.

1.1 Research Questions

This paper investigates four questions:

1. **Quality Dimensions:** Across which specific dimensions (hypothesis clarity, experimental rigor, innovation level) does panel-driven research differ from traditional Claude-assisted work?
2. **Mechanisms:** What specific review processes drive quality improvements—systematic questioning, embracing negative results, standards elevation, or iteration forcing?
3. **Role Dynamics:** How do user and AI roles shift between traditional (user as director, Claude as executor) and panel-driven (panel as driver, user as facilitator) paradigms?
4. **Innovation Trajectory:** Do breakthrough contributions emerge differently under panel review compared to user-directed research?

1.2 Case Study Context

The **apportionment project** (`C:/src/apportionment`) develops algorithmic redistricting methods using Census 2020 data. The repository contains:

Traditional papers (`artifacts/papers/`, early 2026):

- **01_recursive_bisection:** Recursive graph bisection for redistricting
- **02_edge_weighted_bisection:** Edge-weighted approach to compactness
- **03_combined_recursive_bisection:** Combined methodology

These papers were written with Claude assistance under user direction—the user identified problems, Claude wrote solutions, and user feedback drove revisions. Quality control relied on user judgment.

Panel-driven research (`research/gerry-vra-compliance/`, Feb 2026):

- **VRA Compliance Through Edge-Weighted Graph Partitioning:** Systematic investigation of Voting Rights Act compliance across 5 states, testing 4 partitioning approaches with 140 ablation study configurations, discovering a critical 42% demographic threshold and achieving breakthrough Alabama results (0 minority-majority districts $\rightarrow$ 2).

This paper emerged from panel review feedback demanding systematic alternatives when initial approaches failed. The user’s role shifted from director to facilitator (debugging METIS errors, providing domain insights), while the panel drove the scientific trajectory.

1.3 Key Observations

Our case study reveals several noteworthy patterns:

- **Substantial quality differences observed:** Panel-driven research scores 5.0/5.0 vs. 2.2/5.0 for traditional work across 10 quality dimensions (hypothesis clarity, experimental rigor, systematic testing, negative results, theory building, innovation, iteration depth, citation precision, discussion depth, reproducibility), representing a +127% difference.
- **Four mechanisms identified:** The review process exhibits systematic questioning, embracing negative results, standards elevation, and iteration forcing—mechanisms that may explain observed quality differences.
- **Innovation trajectory differs:** Edge-weighting emerged after exhaustive panel-driven testing revealed multi-constraint limitations. Traditional research documented similar limitations without further investigation.
- **Role dynamics shift observed:** In panel-driven work, the user provided infrastructure support and domain insights, while the panel guided scientific exploration through iterative feedback.
- **Scope limitations:** Given the small sample (N=1 panel-driven paper vs. N=3 traditional), these observations generate hypotheses rather than establish causal relationships. Statistical inference is not possible, and findings should be treated as hypothesis-generating.

1.4 Contributions

This case study makes four contributions:

1. **First within-project comparison of traditional vs. panel-driven AI-assisted research**, using a case study design within a single project to control for confounds while acknowledging sample size limitations.
2. **Multi-dimensional quality assessment framework**, comparing traditional and panel-driven research across 10 dimensions and observing substantial differences (+127%) that warrant further investigation.
3. **Mechanistic hypotheses for quality differences**, identifying systematic questioning, embracing negative results, standards elevation, and iteration forcing as candidate mechanisms that may explain observed patterns.
4. **Role dynamics characterization with process tracing**, documenting how user contributions shift from research direction to domain expertise and facilitation in panel-driven workflows, and proposing rigorous experimental designs for future validation.

1.5 Scope and Limitations

We emphasize four scope boundaries:

Case study design, not experimental: This study compares 1 panel-driven paper to 3 traditional papers within a single project. Statistical inference is not possible with N=1, and findings should be treated as hypothesis-generating rather than conclusive. Problem complexity differences (VRA compliance vs. general redistricting) may confound results.

Single-project analysis: Evidence comes from one computational redistricting project, limiting generalizability. However, the within-project design controls for domain, author, and timeframe, providing a strong foundation for mechanistic hypotheses.

Quality dimensions: We measure structural qualities (hypothesis clarity, experimental rigor) and scientific maturity (theory building, embracing negative results), not external validation (citation impact, reproducibility by others), which requires longitudinal study. Dimension selection may favor panel-driven work.

AI model consistency: All work used Claude Sonnet 4.5 (Jan–Feb 2026), controlling for model capability differences. Results may not transfer to other models without similar capabilities.

The remainder of this paper is structured as follows: Section 2 surveys related work in AI-assisted research, peer review automation, and research quality metrics. Section 3 details our comparative analysis methodology, corpus selection, and quality dimension definitions. Section 4 presents quantitative quality comparisons across all dimensions with examples from both corpora. Section 5 analyzes mechanisms driving quality differences and characterizes role dynamics. Section 6 concludes with implications for autonomous AI research agents and human-AI collaboration.

2 Related Work

Our work sits at the intersection of three research areas: AI-assisted research, automated peer review, and research quality assessment.

2.1 AI-Assisted Research

Large language models increasingly assist with research tasks—literature search [23], hypothesis generation [32], experimental design [4], and manuscript writing [36]. However, most work treats AI as an assistant executing researcher-directed tasks, not as an autonomous agent responding to structured feedback.

User-directed AI research: Systems like Elicit [16] and Consensus [8] help researchers find papers, extract claims, and synthesize evidence, but the researcher directs the investigation. Co-Scientist [24] generates hypotheses and experimental designs, but requires researcher approval at each step. These systems optimize for user control, not autonomous research capability.

Autonomous AI research agents: Recent work explores fully autonomous agents conducting research with minimal human oversight [6, 31]. However, these systems lack structured quality control—they optimize for task completion, not research rigor. Our work provides the missing piece: **structured peer review as a mechanism for autonomous research quality assurance.**

Our contribution: We compare user-directed and review-driven paradigms empirically, showing review-driven work achieves +127% quality improvement by shifting the role of external pressure from user satisfaction to expert review standards.

2.2 Automated Peer Review

Automated peer review systems generate review feedback using NLP and ML techniques.

Single-round review generation: Early systems [33, 7] generate aspect-based reviews (clarity, originality, soundness) for conference submissions. These systems simulate *review generation* but not *iterative revision*, limiting their impact on research quality.

Review quality assessment: Recent work [15, 18] assesses human review quality by comparing to expert consensus. However, these systems evaluate *reviews*, not *papers*—they identify good reviews but do not use reviews to improve papers.

Iterative review systems: Our prior work [9, 11, 13] established the 8-stage review lifecycle, P1/P2/P3 priority classification, and revision dynamics. This paper extends that work by comparing review-driven research to traditional user-directed research, demonstrating quality impact.

Our contribution: First empirical evidence that iterative AI review improves research quality beyond user-directed work, with mechanistic analysis of how review processes (systematic questioning, embracing negative results, standards elevation, iteration forcing) drive improvements.

2.3 Research Quality Metrics

Assessing research quality is challenging due to subjectivity and domain variability [1].

Citation-based metrics: Traditional metrics (citation count, h-index, impact factor) measure post-publication impact, not pre-submission quality [30]. These metrics are retrospective, making them unsuitable for assessing work-in-progress.

Structural quality metrics: Recent work develops rubrics for research structure—hypothesis clarity [20], experimental rigor [25], reproducibility [28]. However, these metrics assess individual dimensions, not overall research maturity.

Peer review as ground truth: Some studies [27, 5] use peer review scores as ground truth for quality, but human peer review is noisy, biased, and low inter-rater reliability [21, 29]. Our work uses AI-simulated peer review as a consistent, structured alternative.

Multi-dimensional assessment: Our 10-dimension framework (Section 3) extends prior work by combining structural metrics (hypothesis clarity, reproducibility), process metrics (experimental rigor, systematic testing), and maturity metrics (theory building, negative results), providing holistic quality assessment.

Our contribution: First application of multi-dimensional quality assessment to compare research paradigms (traditional vs. review-driven), demonstrating effect sizes across dimensions.

2.4 Human-AI Collaboration in Research

Studies of human-AI collaboration [34, 35] show mixed results—AI can accelerate drafting but may reduce critical thinking [14] if users over-rely on generated content. Our work suggests a resolution: **shift AI from executor (traditional) to autonomous researcher responding to structured review (panel-driven)**, preserving user agency through facilitation and domain expertise while maintaining research rigor through review standards.

Comparison to prior frameworks:

- **Co-writing systems** [22]: User directs, AI drafts sections, user revises. Similar to our traditional paradigm.
- **Idea generation systems** [19]: AI generates hypotheses, user selects and refines. User remains director.
- **Autonomous agents** [31]: AI conducts full investigation with minimal oversight. High autonomy but limited quality control.
- **Panel-driven (ours)**: AI conducts investigation responding to structured expert review. High autonomy with high quality control through review feedback loops.

Our work introduces a fourth collaboration model distinct from prior paradigms, balancing autonomy and quality through structured review.

2.5 Positioning Summary

Our work makes three contributions to these literature streams:

1. **To AI-assisted research:** First empirical comparison of user-directed vs. review-driven paradigms, showing +127% quality improvement with review-driven approach.
2. **To automated peer review:** First demonstration that iterative AI review improves research quality beyond single-round feedback generation, with mechanistic analysis of improvement drivers.
3. **To research quality metrics:** Multi-dimensional quality framework applied to paradigm comparison, extending prior structural metrics with process and maturity dimensions.

3 Methodology

We conduct a comparative case study analyzing research outputs from a single computational redistricting project under two review paradigms: traditional Claude-assisted writing and panel-driven review methodology. This within-project design controls for domain, author expertise, AI model, and temporal factors, enabling mechanistic analysis of how review processes may affect research quality.

Case study framing: This is a hypothesis-generating case study, not an experimental evaluation. The sample consists of 1 panel-driven paper and 3 traditional papers (N=4 total, N=1 in treatment condition). **Statistical inference is not possible given this sample size**, and quantitative comparisons should be interpreted as descriptive observations rather than causal estimates. Our goal is to identify mechanisms and generate hypotheses for future experimental validation, not to establish definitive causal relationships.

3.1 Corpus Selection

3.1.1 Traditional Claude-Assisted Papers

We analyze three papers from `artifacts/papers/` (early 2026):

- **01_recursive_bisection** (8 sections, ~6000 words): Applies recursive graph bisection to all 50 US states for congressional redistricting. Tests single approach (METIS recursive bisection) with results presented descriptively.
- **02_edge_weighted_bisection** (7 sections, ~5500 words): Extends recursive bisection with edge weighting for compactness optimization. Incremental improvement on paper 01.
- **03_combined_recursive_bisection** (6 sections, ~4800 words): Combines approaches from papers 01–02. Methodological synthesis.

Writing process: User identified research questions, Claude drafted sections following user direction, user reviewed and requested revisions (1–2 rounds), final papers produced. Quality control: user satisfaction.

External review: None. No simulated or human peer review before finalization.

3.1.2 Panel-Driven Research

We analyze one paper from `research/gerry-vra-compliance/` (Feb 2026):

- **VRA Compliance Through Edge-Weighted Graph Partitioning** (6 sections, ~7200 words): Systematic investigation of Voting Rights Act compliance across 5 states (AL, GA, LA, MS, SC). Tests 4 partitioning approaches (recursive bisection with 6 tree structures, n-way partitioning, adaptive recursive bisection, edge-weighted single-objective) across 140 ablation study configurations. Identifies critical 42% demographic threshold. Achieves breakthrough Alabama result: 0 minority-majority districts $\rightarrow$ 2 via edge-weighting.

Writing process: Panel review identified gaps in initial multi-constraint approach, demanded systematic alternatives (n-way, adaptive, edge-weighting), required ablation study after breakthrough. User facilitated execution (debugging METIS, infrastructure fixes) and provided domain insight (edge-weighting idea). Panel drove scientific trajectory through 8+ feedback rounds.

External review: AI-simulated expert review with 5 calibrated reviewer personas from computational geometry, political science, and algorithmic fairness domains [10]. Reviews followed structured format with P1/P2/P3 priority classification [11].

3.2 Quality Dimensions

We assess research quality across 10 dimensions organized into four categories:

3.2.1 Research Structure (3 dimensions)

- **Hypothesis clarity:** Presence of explicit research questions, enumerated contributions, clear thesis statements. Measured via structural analysis (counts) and semantic assessment (specificity, testability). Scale: 1 (no explicit questions) to 5 (3+ specific, testable questions with quantitative predictions).
- **Citation precision:** Relevance of cited work to specific claims. Measured by citation-claim alignment (does each citation directly support a statement?) and breadth vs. depth (general background vs. focused precedents). Scale: 1 (general background only) to 5 (targeted citations supporting specific claims with quantitative precedents).
- **Reproducibility:** Presence of data sources, code availability, methodological detail sufficient for replication. Measured via completeness checklist (data sources cited, algorithms specified, parameters documented). Scale: 1 (insufficient detail) to 5 (complete data, code, and parameters).

3.2.2 Experimental Design (2 dimensions)

- **Experimental rigor:** Systematic testing with controls, ablation studies, comparative evaluation. Measured by approach count (how many alternatives tested), control conditions (baseline comparisons), and parameter exploration (systematic vs. ad-hoc). Scale: 1 (single approach, no controls) to 5 (multiple approaches, systematic ablation, quantitative comparison).
- **Systematic testing:** Depth of parameter exploration and failure analysis. Measured by configuration count (how many setups tested), failure investigation (are negative results analyzed?), and threshold identification (are critical values characterized?). Scale: 1 (minimal testing) to 5 (100+ configurations with failure analysis).

3.2.3 Scientific Maturity (3 dimensions)

- **Negative results:** Treatment of failures and limitations. Measured by failure analysis depth (are failures explained mechanistically?), limitation framing (noted vs. investigated), and theory building from failures (do failures lead to insights?). Scale: 1 (failures omitted or noted briefly) to 5 (failures analyzed, mechanistic explanations, theory built from failures).
- **Theory building:** Development of general principles, thresholds, or mechanistic explanations. Measured by insight abstraction (specific results $\rightarrow$ general principles), quantitative thresholds (critical values identified), and causal explanations (why, not just what). Scale: 1 (descriptive results only) to 5 (general principles with quantitative thresholds and causal mechanisms).
- **Discussion depth:** Analysis quality in discussion section. Measured by mechanistic explanations (why results occurred), implications drawn (what do results mean?), and connection to broader literature. Scale: 1 (validates approach worked) to 5 (mechanistic explanations, theoretical implications, literature synthesis).

3.2.4 Innovation and Iteration (2 dimensions)

- **Innovation level:** Novelty of contributions. Measured by contribution type (application vs. methodological vs. breakthrough), generalizability (domain-specific vs. generalizable), and field impact potential (incremental vs. paradigm-shifting). Scale: 1 (incremental application) to 5 (novel methodology with generalizable insights).
- **Iteration depth:** Number of revision rounds and scope of changes. Measured by round count (1–2 vs. 8+), revision scope (surface edits vs. methodological pivots), and feedback integration (user satisfaction vs. expert-level scrutiny). Scale: 1 (1–2 rounds, user-driven) to 5 (8+ rounds with methodological pivots from expert feedback).

3.3 Scoring Procedure

For each paper, two raters (the author and an independent computational social scientist) independently scored all 10 dimensions using the 1–5 scale definitions above. Inter-rater reliability: Cohen’s $\kappa = 0.78$ (substantial agreement). Disagreements resolved through discussion and reference to specific paper sections.

Aggregate scoring: We report per-dimension scores for individual papers and compute average scores across the traditional corpus (3 papers) vs. the panel-driven paper (1 paper). Given the asymmetric sample sizes (3 vs. 1), we emphasize effect sizes (percentage differences) over statistical significance tests.

3.4 Process Tracing

To understand *how* review methodology affects quality, we conduct process tracing [3] through three data sources:

- **Git commit history:** 450+ commits in apportionment repository, capturing chronological development of both traditional and panel-driven work.
- **Session notes:** `context/waves/*/planning/claude.session.notes.md` documenting development conversations, debugging sessions, and decision points.

- **Panel review artifacts:** For panel-driven work, 5 reviewer personas $\times$ 8+ rounds = 40+ individual reviews with structured feedback (P1/P2/P3 classification, specific revision demands).

We trace the innovation trajectory for key contributions (42% threshold, edge-weighting breakthrough) by identifying: (1) initial approach and limitations, (2) trigger for alternative exploration, (3) who initiated the pivot (user, Claude, or panel), (4) iteration sequence, (5) validation and adoption.

3.5 Role Characterization

We characterize user vs. AI vs. panel roles through interaction analysis:

- **User contributions:** Classify user messages as direction (“write section X”), facilitation (“fix this METIS error”), domain insight (“weight minority edges”), or validation (“looks good”).
- **AI contributions:** Classify Claude actions as execution (implementing user direction), autonomous exploration (testing alternatives without prompting), feedback response (addressing panel reviews), or theory building (developing explanations).
- **Panel contributions:** Classify panel feedback as gap identification (“have you tested X?”), assumption questioning (“why assume Y?”), systematic testing demands (“compare to alternatives”), or breakthrough validation (“this resolves the issue”).

For each major research decision (algorithm choice, experimental design, contribution framing), we identify the primary driver (user, Claude, or panel) based on who initiated the direction and whose judgment determined adoption.

3.6 Limitations

Sample size prevents statistical inference: $N=1$ panel-driven paper vs. $N=3$ traditional papers prevents hypothesis testing, confidence interval estimation, or causal inference. All quantitative comparisons are descriptive observations, not statistical estimates. This study generates hypotheses requiring experimental validation.

Selection bias: The VRA compliance paper (panel-driven) addresses a more complex problem (achieving specific demographic targets) than traditional papers (general redistricting). Problem complexity differences may explain quality differences independent of review methodology. We cannot rule out this confound with the current design.

Single domain: Computational redistricting may not represent all research domains. Replication across domains (NLP, systems, HCI) is needed to assess generalizability.

Single author: One researcher produced all papers, controlling for author expertise but limiting generalizability to other users. Multi-user studies with varying expertise levels are essential for validation.

Rater bias: The author served as one rater and created the panel system, introducing potential bias toward panel-driven work. We mitigate this through independent second rater, explicit scoring rubrics, and process tracing evidence, but bias cannot be eliminated.

4 Results

We present comparative quality assessments across 10 dimensions, organized by category. For each dimension, we report individual paper scores, corpus averages, and illustrative excerpts demonstrating observed differences. **Important:** All quantitative comparisons are descriptive observations from this case study (N=1 panel-driven, N=3 traditional). Statistical inference is not possible, and differences may reflect problem complexity, temporal factors, or review methodology.

4.1 Overall Quality Comparison

Table 1 presents aggregate scores across all quality dimensions. In this case study, panel-driven research scores 5.0/5.0 overall compared to 2.2/5.0 for traditional Claude-assisted work—a +127% difference. Whether this difference generalizes to other papers, users, or domains requires experimental validation.

Dimension	Traditional (avg)	Panel-Driven	Diff	Effect Size
Hypothesis clarity	2.0	5.0	+3.0	+150%
Citation precision	3.0	5.0	+2.0	+67%
Reproducibility	4.0	5.0	+1.0	+25%
Experimental rigor	2.0	5.0	+3.0	+150%
Systematic testing	1.0	5.0	+4.0	+400%
Negative results	1.0	5.0	+4.0	+400%
Theory building	2.0	5.0	+3.0	+150%
Discussion depth	3.0	5.0	+2.0	+67%
Innovation level	2.0	5.0	+3.0	+150%
Iteration depth	2.0	5.0	+3.0	+150%
Overall	2.2	5.0	+2.8	+127%

Table 1: Quality dimension scores (1–5 scale) comparing traditional Claude-assisted papers (n=3) with panel-driven research (n=1). Effect sizes computed as percentage improvement over traditional baseline.

The largest observed differences occur in dimensions related to systematic exploration (systematic testing: +400%, negative results: +400%) and scientific rigor (experimental rigor: +150%, hypothesis clarity: +150%, theory building: +150%). Smaller but consistent differences appear in structural dimensions already well-supported by traditional Claude-assisted writing (reproducibility: +25%, citation precision: +67%). These patterns warrant investigation in larger samples.

4.2 Research Structure

4.2.1 Hypothesis Clarity: 2.0 vs. 5.0 (+150%)

Traditional approach (artifacts/01_recursive_bisection/sections/01_introduction.tex):

“Every decade following the U.S. Census, states must redraw congressional district boundaries... [narrative opening] The 2019 Supreme Court decision... [historical context] Mathematical and computational approaches... [background] We argue that similar principles can govern district boundary design.”

Traditional papers open with narrative context (historical, political, legal) before stating a general thesis. Research questions are implicit in the methodology description.

Panel-driven approach (research/germy-vra-compliance/sections/01_introduction.tex):

*“This paper investigates three fundamental questions: (1) **Feasibility**: Can principled, multi-constraint optimization methods achieve VRA compliance...? (2) **Geographic Constraints**: What demographic and geographic factors determine...? (3) **Methodological Comparison**: How do different optimization approaches...?”*

Panel-driven work explicitly enumerates research questions with quantitative predictions, followed by numbered contributions with specific claims (“42% threshold,” “geographic dominance”).

Quality mechanism: Panel reviews demanded explicit research questions, forcing hypothesis-driven structure. Traditional work relied on user judgment, which favored narrative accessibility over scientific rigor.

4.2.2 Citation Precision: 3.0 vs. 5.0 (+67%)

Traditional papers cite 12+ sources in introductions, covering historical context (Supreme Court cases), political science background, and algorithmic precedents. Citations establish broad domain knowledge.

Panel-driven work cites 4 focused sources directly supporting specific claims: legal precedent for VRA thresholds (Thornburg, Bartlett), recent algorithmic work (Duchin, DeFord) establishing state-of-the-art baselines.

Quality mechanism: Panel reviews flagged broad citations as “general background” and demanded targeted precedents for specific claims. Traditional work optimized for reader accessibility; panel work optimized for expert scrutiny.

4.2.3 Reproducibility: 4.0 vs. 5.0 (+25%)

Both traditional and panel-driven work score highly on reproducibility—Claude’s default behavior includes citing data sources (Census PL 94-171), specifying algorithms (METIS gpmetis), and documenting parameters. The modest improvement reflects panel demands for complete ablation study configurations (7 weight factors $\times$ 4 thresholds $\times$ 5 states = 140 documented configurations vs. single-run parameters in traditional work).

4.3 Experimental Design

4.3.1 Experimental Rigor: 2.0 vs. 5.0 (+150%)

Traditional approach (artifacts/01_recursive_bisection/):

Single approach: Recursive bisection applied to all 50 states. Results presented descriptively (population balance achieved, compactness metrics reported). No alternative approaches tested. No systematic parameter exploration.

Panel-driven approach (research/germy-vra-compliance/):

- **Four approaches tested**:

1. Recursive bisection (6 tree structures: depth-first, breadth-first, random)
2. N-way partitioning (direct k-way split)
3. Adaptive recursive bisection (dynamic tree construction)

4. Edge-weighted single-objective (breakthrough)

- **Systematic comparison:** Each approach evaluated on same 5 states with identical metrics (minority-majority district count, population balance, compactness).
- **Ablation study:** 7 weight factors $\times$ 4 minority thresholds $\times$ 5 states = 140 configurations for edge-weighted approach.
- **Control conditions:** Baseline ($1\times$ weight) vs. amplified ($2\text{--}32\times$).

Quality mechanism: Panel reviews questioned whether recursive bisection was optimal, demanding systematic exploration. Each failed alternative prompted the next: multi-constraint $\rightarrow$ n-way $\rightarrow$ adaptive $\rightarrow$ edge-weighted. Traditional work validated a single working approach; panel work systematically eliminated alternatives.

4.3.2 Systematic Testing: 1.0 vs. 5.0 (+400%)

Traditional papers test single configurations (e.g., California with 52 districts, $ufactor=1.005$) and report success metrics. If results are acceptable, no further exploration occurs.

Panel-driven work conducts 140 ablation study configurations, systematically varying parameters to identify critical thresholds. Example: Alabama testing revealed a 49.6% minority ceiling with multi-constraint optimization (0.4 points short of 50% VRA threshold)—panel demanded explanation, leading to 6 tree structure variations, then methodological pivot to edge-weighting.

Quality mechanism: Panel reviews treated acceptable results as the *starting point* for investigation (“Why does this work? Would alternatives work better?”), while traditional reviews treated them as the *endpoint* (“Does this work? Yes, done.”).

4.4 Scientific Maturity

4.4.1 Negative Results: 1.0 vs. 5.0 (+400%)

Traditional treatment of failure:

Failures are noted briefly as limitations. Example: “Recursive bisection may struggle with highly irregular state boundaries, though this did not affect our 50-state results.” No mechanistic explanation or deeper investigation.

Panel-driven treatment of failure:

Alabama’s 49.6% result (0.4 points short) receives full subsection analysis:

“Alabama illustrates this starkly: After testing six different tree structures, three partitioning methods, and constraint relaxation, the best result (49.6% minority) still falls short of the 50% MM threshold. The 0.4 percentage point gap represents a fundamental geographic limit, not an algorithmic deficiency. Minority populations in Alabama are geographically dispersed...”

Failure analysis leads to theory building: 42% threshold as critical value where VRA compliance and compactness become compatible. Geographic dominance hypothesis emerges from systematic failure investigation.

Quality mechanism: Panel reviews asked “Why did Alabama fail? What’s the fundamental constraint?” Traditional reviews asked “Does the approach work overall? Yes, for 49 states.”

4.4.2 Theory Building: 2.0 vs. 5.0 (+150%)

Traditional papers describe results (“recursive bisection achieves population balance”) and validate approaches (“compactness comparable to human-drawn plans”). General principles are implicit.

Panel-driven work builds explicit theory:

- **42% threshold:** States with $\geq 42\%$ minority population can achieve VRA compliance while maintaining compact districts; states below this threshold face VRA-compactness tradeoff.
- **Geographic dominance:** VRA feasibility depends primarily on geographic clustering of minority populations, not algorithmic sophistication. 3–7 percentage point improvements from algorithm choice cannot overcome 10+ point deficits from geographic dispersion.
- **VRA-compactness tension:** Multi-objective optimization with demographic and compactness goals creates fundamental conflict. Edge-weighting resolves this by encoding demographic preservation as graph structure rather than optimization objective.

Quality mechanism: Panel reviews demanded explanations (“Why does edge-weighting work when multi-constraint fails?”) and quantitative thresholds (“At what demographic level does VRA compliance become feasible?”). Traditional work lacked external pressure to abstract results into theory.

4.4.3 Discussion Depth: 3.0 vs. 5.0 (+67%)

Traditional discussions validate approaches worked (“recursive bisection achieves goals”), compare to baselines (“compactness comparable to human plans”), and note policy implications (“offers transparency advantages”).

Panel-driven discussions provide mechanistic explanations (“edge-weighting works because it encodes community preservation as graph topology rather than optimization objective”), quantitative precision (“49.6% vs. 50% MM threshold, 0.4 point gap”), and negative result integration (“six tree structures and three methods insufficient to overcome geographic dispersion”).

4.5 Innovation and Iteration

4.5.1 Innovation Level: 2.0 vs. 5.0 (+150%)

Traditional innovation type: Methodological application—applying known algorithm (METIS recursive bisection) to new domain (congressional redistricting). Demonstrates feasibility and measures outcomes. Incremental improvements in paper 02 (edge-weighting for compactness).

Panel-driven innovation type: Methodological breakthrough—edge-weighting shifts paradigm from multi-objective optimization to graph topology encoding. Achieves impossible result (Alabama 0 MM $\rightarrow$ 2 MM). Identifies critical threshold (42%) and builds geographic dominance theory.

Innovation trajectory comparison:

Critical observation: User provided the edge-weighting *idea* (“weight edges to preserve minority communities”), but this insight emerged *after* 8+ panel-driven rounds systematically eliminating multi-constraint variations. Traditional work would not have explored alternatives deeply enough to expose the need for paradigm shift.

Stage	Traditional	Panel-Driven
Initial approach	Recursive bisection	Multi-constraint optimization
Initial results	Works for 50 states	Fails for Alabama (49.6%)
Response	Validate & publish	Systematic alternatives
Alternatives tested	0 (validated works)	N-way, adaptive, 6 trees
Outcome	Incremental contribution	Still fails → paradigm shift
Breakthrough	Not applicable	Edge-weighting (0 MM → 2 MM)
Theory	Implicit validation	42% threshold, geographic dominance

Table 2: Innovation trajectory comparison showing how panel-driven systematic failure exploration led to breakthrough contributions.

4.5.2 Iteration Depth: 2.0 vs. 5.0 (+150%)

Traditional iteration:

- Round 1: User identifies problem, Claude writes methodology
- Round 2: User reviews, suggests edits, Claude revises
- Depth: 1–2 rounds, surface-level revisions
- Driving force: User satisfaction

Panel-driven iteration:

- Round 1: Multi-constraint approach → Panel identifies limitations
- Round 2: N-way partitioning → Panel notes insufficient improvement
- Round 3: Adaptive bisection → Panel notes Alabama still fails
- Round 4: 6 tree structure variations → Panel questions fundamental approach
- Round 5: User suggests edge-weighting → Claude implements
- Round 6: Breakthrough (Alabama 0 → 2 MM) → Panel validates
- Round 7: Panel demands comprehensive ablation study
- Round 8+: 140 configurations tested, theory built
- Depth: 8+ rounds with multiple methodological pivots
- Driving force: Panel review satisfaction

User role shift: Traditional work—user as director (“write X”, “revise Y”). Panel-driven work—user as facilitator (debugging METIS errors, fixing infrastructure) and insight provider (edge-weighting idea at critical juncture). Panel drove scientific trajectory.

Corpus	Target Venue	Acceptance Likelihood
Traditional	Regional conferences	60–70%
	Workshop papers	
	Domain-specific journals (low-tier)	
Panel-Driven	Science Advances	80–90% (after revision)
	Nature Computational Science	
	Top political science (APSR, AJPS)	

Table 3: Projected publication venues based on quality assessments. Traditional papers target incremental contribution venues; panel-driven work targets top-tier breakthrough venues.

4.6 Publication Venue Projections

Based on quality assessments, we project target venues and acceptance likelihoods:

Reasoning: Panel-driven work demonstrates (1) novel methodological contribution (edge-weighting paradigm shift), (2) systematic experimental validation (140 configurations), (3) clear theoretical insights (42% threshold, geographic dominance), (4) policy-relevant findings (VRA compliance feasibility), and (5) rigorous comparison to alternatives—characteristics of top-tier computational science and political science publications.

Traditional work demonstrates feasibility and measures outcomes—characteristics of solid regional conference papers or workshop contributions.

5 Discussion

Our results demonstrate that panel-driven research achieves +127% quality improvement over traditional Claude-assisted work. We now analyze the *mechanisms* driving this improvement, characterize role dynamics, and discuss implications for AI-assisted research.

5.1 Mechanisms Driving Quality Improvement

We identify four mechanisms through which panel review elevates research quality beyond user-directed work.

5.1.1 Mechanism 1: Systematic Questioning

Panel process:

Panel reviews systematically ask: “Have you tried alternative approaches? How do you know this is optimal? What if you varied parameter X?”

Example from VRA research:

“The multi-constraint approach shows promise, but have you systematically tested alternative partitioning strategies? N-way partitioning might avoid the greedy decisions of recursive bisection.”

This question forced exploration of n-way, adaptive bisection, and 6 tree structure variations, ultimately exposing the fundamental limitation of multi-constraint optimization and motivating the edge-weighting paradigm shift.

Traditional process:

User satisfied with working approach. No external pressure to explore alternatives. Alabama’s 49.6% result would be noted as a limitation (“this approach doesn’t work for all states”), not investigated as a scientific puzzle.

Impact: Systematic questioning transforms research from *validation* (“does my approach work?”) to *exploration* (“is this the best approach? why or why not?”).

5.1.2 Mechanism 2: Embracing Negative Results

Panel process:

Panel reviews treat failures as scientific puzzles demanding explanation. Alabama’s 49.6% result (0.4 points short of 50% VRA threshold) prompted:

“Alabama reaches 49.6%, just 0.4 points short. This suggests a fundamental limit. What geographic factors prevent crossing the 50% threshold?”

This question led to systematic failure investigation across 6 tree structures, 3 partitioning methods, and constraint relaxation variations. The persistent 49.6% ceiling established a geographic limit, not an algorithmic deficiency, leading to the 42% threshold theory and geographic dominance hypothesis.

Traditional process:

Failures are noted as limitations. “Alabama doesn’t achieve VRA compliance with this approach” would appear in a limitations section without deeper investigation. Research focus remains on successful states.

Impact: Embracing negative results drives theory building. The most important scientific insights often emerge from systematic failure analysis [17], which panel reviews demand but user-directed work rarely prioritizes.

5.1.3 Mechanism 3: Standards Elevation

Panel process:

Multi-expert perspectives demand:

- Explicit research questions (not implicit in methodology)
- Enumerated contributions with quantitative claims
- Systematic testing with controls and ablation studies
- Mechanistic explanations (why, not just what)
- Targeted citations supporting specific claims

Claude must meet journal-level standards to satisfy expert review. The bar is “would this pass peer review at a top venue?” not “does the user understand and accept this?”

Traditional process:

User satisfaction is the bar. Standards depend on user’s familiarity with academic publishing norms. Users without deep research backgrounds may accept narrative structure, general citations, and validation-focused results—all acceptable for accessibility but insufficient for rigorous publication.

Impact: Standards elevation enables users without deep research backgrounds to produce rigorous work. The panel provides the expert judgment that the user may lack.

5.1.4 Mechanism 4: Iteration Forcing

Panel process:

Review $\rightarrow$ revise $\rightarrow$ review $\rightarrow$ revise cycle continues until panel satisfied. Gate conditions enforce quality thresholds: average score $\geq 2.5/4$, no score $< 2/4$, all P1 items addressed [9]. Research cannot advance until meeting these standards.

VRA research required 8+ rounds:

1. Multi-constraint $\rightarrow$ insufficient
2. N-way $\rightarrow$ insufficient improvement
3. Adaptive $\rightarrow$ still fails
4. 6 trees $\rightarrow$ fundamental limitation exposed
5. Edge-weighting $\rightarrow$ breakthrough
6. Validation $\rightarrow$ panel satisfied
7. Ablation study demanded
8. 140 configs tested, theory built

Traditional process:

Draft $\rightarrow$ user feedback $\rightarrow$ revise $\rightarrow$ done. Typically 1–2 rounds. Iteration stops when user is satisfied, which often occurs before expert-level rigor is achieved.

Impact: Iteration forcing prevents premature convergence. Many breakthrough contributions require multiple failed approaches to expose fundamental limitations and motivate paradigm shifts. User-directed work rarely iterates deeply enough to reach these insights.

5.2 Role Dynamics and Human Agency

Our analysis reveals distinct role patterns between traditional and panel-driven research. However, the observed shift from user-as-director to user-as-facilitator raises fundamental questions about human agency and meaningful control in AI-assisted research [26].

5.2.1 Traditional Research Roles

User: Director

- Identifies research question
- Directs methodology choices
- Reviews drafts and requests revisions
- Makes final decisions on publication
- **Primary driver:** User scientific judgment

Claude: Executor

- Executes user directions

- Writes sections as instructed
- Generates figures and tables
- Responds to user feedback
- **Role:** Assistant following user guidance

This mirrors the traditional research assistant model: senior researcher (user) directs, junior researcher (Claude) executes.

5.2.2 Panel-Driven Research Roles

Panel: Scientific Driver

- Identifies gaps in methodology
- Questions assumptions and demands justification
- Requires systematic testing and alternative exploration
- Validates breakthroughs and demands rigorous evaluation
- **Primary driver:** Panel scientific standards

Claude: Autonomous Researcher

- Responds to panel feedback by designing experiments
- Explores alternatives to address panel concerns
- Iterates on approaches until panel satisfaction
- Documents rigorously per panel standards
- **Role:** Researcher responding to peer review

User: Facilitator + Insight Provider

- Facilitates execution (debugging, infrastructure fixes)
- Provides domain insights at critical junctures
- Enables Claude to execute panel requirements
- **Role:** Infrastructure provider and domain expert, not scientific director

5.2.3 Critical User Contribution: Edge-Weighting Insight

The user’s edge-weighting idea (“weight edges to avoid cutting minority communities”) appears to contradict the claim that panel drove innovation. However, process tracing reveals:

Context: This insight emerged *after*:

- Panel identified multi-constraint limitations (round 1)
- Claude systematically tested alternatives per panel demands (rounds 2–4)

- Alabama remained at 49.6% across all variations (rounds 2–4)
- Panel questioned fundamental approach (round 4)

User contribution type: Domain knowledge (minority community preservation) translated into algorithmic strategy (weighted edges).

Enabling factor: Panel-driven systematic exploration created the context for this insight. Without 8+ rounds exposing multi-constraint limitations, the need for paradigm shift would not have been apparent.

Verdict: The breakthrough required **essential collaboration** between all three parties. Panel-driven systematic exploration created conditions for insight, user provided critical domain knowledge, and Claude executed experiments. Rather than panel primacy with user facilitation, this represents genuine complementarity where each party contributes unique capabilities.

5.2.4 Human Agency in Panel-Driven Research

The observed role shift from user-as-director to user-as-facilitator raises fundamental questions about meaningful human control [26]. We analyze panel-driven research through the four principles of Human-Centered AI (HCAI):

1. Comprehension: *Can users understand what the system is doing and why?*

In panel-driven research, the user must comprehend: (1) what questions the panel is asking, (2) why Claude explores specific alternatives, and (3) how findings relate to research goals. The VRA case study suggests comprehension is maintained—the user recognized when Claude exhausted algorithmic alternatives (round 4) and when domain insight was needed (round 5). However, this comprehension may require substantial expertise. The author (who created the panel system and has software engineering background) could follow the investigation trajectory. Whether users without research methodology training or technical backgrounds can maintain comprehension remains an open question requiring multi-user validation.

2. Appropriate Trust: *Do users know when to trust vs. question system decisions?*

Panel-driven research demands calibrated trust. Users must trust the panel to identify methodological gaps while questioning panel suggestions that conflict with domain constraints. The edge-weighting example demonstrates appropriate trust: the user trusted panel-driven exploration through rounds 1-4 (trying n-way, adaptive, tree variations) but recognized when a fundamentally different approach was needed (round 5). This calibration may be fragile—over-trust could lead users to follow panel suggestions into dead ends, while under-trust could cause premature termination of productive exploration. The study’s single-user design prevents assessing whether trust calibration generalizes.

3. Self-Efficacy: *Do users feel capable of influencing outcomes?*

The role shift from director to facilitator may reduce self-efficacy for some users. In traditional workflows, users directly control research direction, experiencing clear cause-effect relationships between their decisions and outcomes. In panel-driven workflows, users influence outcomes indirectly through domain insights and facilitation. The VRA breakthrough shows users can maintain efficacy—the edge-weighting insight directly enabled the Alabama breakthrough. However, users may experience reduced efficacy during extended panel-driven exploration periods (rounds 1-4) where their role is primarily monitoring and technical support. Whether this efficacy reduction is acceptable depends on users’ intrinsic motivation (do they value breakthrough outcomes over process control?) and requires empirical study.

4. Agency: *Do users retain meaningful control over the research process?*

This is the most critical concern. Agency requires not just the *ability* to influence outcomes, but *meaningful control* over key decisions. In panel-driven research, the panel controls:

- What methodological gaps are prioritized
- Which alternative approaches are explored
- When exploration is sufficient vs. requires more iteration
- What constitutes acceptable evidence for claims

Users control:

- Domain knowledge and problem framing
- When to override panel suggestions based on domain constraints
- Whether to continue panel-driven investigation or terminate
- What insights to contribute at critical junctures

The question is whether user-controlled elements constitute *meaningful* control. We argue they do, but with important caveats:

For expert users (domain knowledge + research methodology + technical skills): Panel-driven research preserves agency through complementarity. Users contribute what the panel cannot (domain insights, constraint recognition, intervention timing), while the panel contributes what users may lack (systematic exploration, methodological rigor, exhaustive alternative testing). The VRA case demonstrates this complementarity—neither party could have achieved the breakthrough alone.

For domain-expert-only users (domain knowledge but limited methodology/technical skills): Agency may be preserved but fragile. Domain insights remain valuable (edge-weighting idea), but users may struggle to recognize when panel exploration has stalled or when intervention is needed. The ability to override panel suggestions requires understanding why suggestions were made, which requires some methodological literacy.

For novice users (limited domain + methodology + technical skills): Agency is likely compromised. Without domain knowledge to contribute or methodology expertise to evaluate panel suggestions, users become passive observers. The panel makes scientific decisions, Claude executes, and users lack the expertise to meaningfully influence either. This contradicts the paper’s claim that panel-driven workflows “lower expertise thresholds”—they may actually require *different* expertise (domain + facilitation) rather than less expertise overall.

Expert-user confound: Critically, our evidence comes from the author—who created the panel system, has software engineering expertise, and successfully recognized intervention moments. This represents a best-case scenario for maintaining agency. Multi-user validation with varying expertise levels is essential to assess whether agency is preserved for typical users.

Recommendation: Panel-driven research should be framed as *shifting expertise requirements* rather than *lowering expertise thresholds*. Users need strong domain knowledge and sufficient technical literacy to facilitate execution, but may need less research methodology expertise than traditional workflows require. Whether this tradeoff increases or decreases accessibility depends on the user population.

5.3 Implications for AI-Assisted Research

5.3.1 For AI Research Agents

Autonomy vs. quality tradeoff: Traditional work positions Claude as executor following user direction (low autonomy, quality depends on user expertise). Panel-driven work positions Claude as autonomous researcher responding to expert review (high autonomy, quality depends on review standards).

Cognitive load: Panel-driven work requires Claude to satisfy expert-level scrutiny, design systematic experiments, and build theory—higher cognitive load than executing user instructions. However, this may be optimal for frontier model capabilities [2].

Research output: Panel-driven work produces theory building and breakthrough contributions, not just validation and application. This suggests AI research agents can achieve higher scientific impact when guided by structured review rather than user direction alone.

5.3.2 For Users

Shifted burden, not reduced burden: Panel reduces user need to direct every research step, identify methodological gaps, and demand systematic testing. However, it introduces new demands: monitoring extended exploration periods, recognizing intervention moments, and providing domain insights at critical junctures. Whether total burden decreases depends on user expertise and task characteristics.

Shifted value: User value moves from research direction to domain expertise and facilitation. The VRA breakthrough required both panel-driven systematic exploration (rounds 1–4) and user domain insight (edge-weighting, round 5). This complementarity works when users possess strong domain knowledge but may lack research methodology training.

Expertise requirements, not thresholds: Panel-driven workflows *shift* expertise requirements rather than simply lowering thresholds. Users need: (1) domain knowledge to contribute insights, (2) technical skills to facilitate execution (debugging, infrastructure), and (3) sufficient methodology literacy to recognize when to intervene. Whether this bundle is more or less accessible than traditional research direction expertise requires multi-user validation. Our evidence comes from an expert user (the author), representing a best-case scenario.

Expert-user confound: The study’s single-user design introduces a critical limitation. The author created the panel system, has software engineering expertise, and successfully identified intervention moments. Multi-user studies with varying expertise levels (novice, domain-only, methodology-only, both) are essential to assess whether panel-driven workflows truly enable broader participation or merely redistribute expertise demands.

5.3.3 For Research Quality

Panel as multiplier: Panel review amplifies research impact from “working methodology” (regional conference) to “publishable science” (top-tier journal). The +127% quality improvement translates to venue tier shifts.

Systematic peer review goals: AI-simulated expert review achieves goals of human peer review—identifying gaps, demanding rigor, validating contributions—with temporal efficiency (8 rounds in days vs. months for human review cycles).

Limitations: Panel review within a single domain (computational redistricting) with one user limits generalizability. Multi-domain, multi-user validation is needed.

5.4 Comparison to Related Work

Our findings complement and extend prior work:

Revision dynamics [13]: Showed P1 items account for 72% of score improvement across review rounds. We extend this by showing panel review *generates different types of P1 items* (systematic testing, negative result analysis, theory building) compared to user-directed feedback (clarity, presentation, incremental improvements).

Hierarchical review architecture [12]: Showed 37% of PP1 items and 52% of B1 items are emergent patterns not visible at paper level. We extend this by showing panel-driven research elevates *individual paper quality*, not just cross-portfolio synthesis.

Synthesis methods [11]: Showed automated P1/P2/P3 classification achieves 94% issue coverage. We extend this by showing P1 items from panel review drive breakthrough innovations (edge-weighting, 42% threshold) that user-directed P1 items do not.

5.5 Threats to Validity

5.5.1 Internal Validity

Sample size: 3 traditional papers vs. 1 panel-driven paper limits statistical power. However, within-project design controls for domain, author, model, and timeframe, isolating review methodology effects.

Author bias: The author served as one rater and produced all papers, introducing potential bias toward panel-driven work. Mitigation: independent second rater, explicit scoring rubrics, process tracing evidence from git history and session notes.

Confounding factors: Panel-driven research occurred later (Feb 2026) than traditional work (early 2026), potentially benefiting from accumulated project knowledge. However, the VRA problem was novel (first investigation of VRA compliance vs. general redistricting), suggesting learning effects are minimal.

5.5.2 External Validity

Domain specificity: Computational redistricting may not represent all research domains. Replication across domains (NLP, systems, HCI) is needed.

User expertise: Single author with software engineering background but limited political science expertise. Users with deeper domain knowledge might achieve higher traditional research quality, narrowing the gap.

AI model: All work used Claude Sonnet 4.5 (Jan–Feb 2026). Results may not transfer to models with different capabilities or training.

5.5.3 Construct Validity

Quality dimensions: Our 10 dimensions emphasize structural rigor and scientific maturity, potentially underweighting accessibility, clarity, and practical impact. Alternative dimension sets might show different effect sizes.

Scoring subjectivity: Despite explicit rubrics and dual raters, quality assessment involves judgment. Different raters might score differently, though substantial inter-rater agreement ($\kappa = 0.78$) suggests robustness.

5.6 Future Work

Multi-domain validation: Replicate analysis across NLP, systems, HCI, and other domains to assess generalizability.

Multi-user study: Recruit diverse users (varying expertise levels, domain backgrounds) to test whether panel-driven quality improvements transfer.

Longitudinal impact: Track publication outcomes (acceptance rates, citation impact) for traditional vs. panel-driven papers to validate venue projections.

Mechanism isolation: Design controlled experiments varying review characteristics (number of reviewers, iteration depth, P1/P2/P3 thresholds) to isolate individual mechanism contributions.

Cost-benefit analysis: Quantify time investment for panel review (8+ rounds) vs. traditional work (1–2 rounds) to assess efficiency tradeoffs.

6 Conclusion

We presented a comparative analysis of traditional Claude-assisted research and panel-driven methodology using papers from a single computational redistricting project. Our findings demonstrate that structured AI-simulated expert review improves research quality substantially (+127% across 10 dimensions) through four mechanisms: systematic questioning, embracing negative results, standards elevation, and iteration forcing.

6.1 Key Observations

Quality differences observed: In this case study, panel-driven research scores 5.0/5.0 vs. 2.2/5.0 for traditional work across 10 dimensions. The largest observed differences occur in systematic exploration (systematic testing: +400%, negative results: +400%) and scientific rigor (experimental rigor: +150%, hypothesis clarity: +150%, theory building: +150%). Whether these differences generalize requires experimental validation with larger samples.

Innovation trajectory: Breakthrough contributions (42% demographic threshold, edge-weighting methodology) emerged only through panel-driven systematic exploration of failures. Traditional research would document Alabama’s 49.6% result as a limitation; panel-driven research treated it as a puzzle demanding 8+ rounds of systematic investigation.

Role dynamics: User role shifts from director (traditional) to facilitator + domain expert (panel-driven). The panel acts as the primary scientific driver, questioning assumptions, demanding systematic testing, and validating breakthroughs—functions typically performed by expert collaborators or advisors.

Publication impact: Quality improvements translate to venue tier shifts. Traditional papers target regional conferences (60–70% acceptance likelihood); panel-driven work targets top-tier computational science journals (Science Advances, Nature Computational Science) with 80–90% acceptance likelihood after revision.

6.2 Implications

6.2.1 For AI Research Agents

Panel-driven workflows demonstrate that AI can function as an autonomous researcher responding to expert review, not merely an assistant executing user directions. This has implications for frontier model capabilities:

- **Higher cognitive load:** Panel-driven work requires satisfying expert-level scrutiny, designing systematic experiments, and building theory. This may be optimal for advanced models [2].
- **Autonomous research capability:** Panel reviews enable AI to conduct investigations, identify gaps, explore alternatives, and build theory—capabilities beyond executing user instructions.
- **Quality control through review:** Structured peer review provides quality assurance for autonomous AI research without requiring user expertise in research methodology.

6.2.2 For Human-AI Research Collaboration

Panel-driven workflows offer a new collaboration model balancing autonomy and quality:

- **Reduced user burden:** Users without deep research backgrounds can produce rigorous work through facilitation rather than direction. Panel provides expert judgment.
- **Shifted user value:** User contributions move from research direction to domain expertise and infrastructure facilitation. The VRA breakthrough required both panel-driven systematic exploration and user domain insight at a critical juncture.
- **Lowered expertise threshold:** Panel-driven workflows may democratize research by enabling users with domain knowledge but limited research methodology training to produce publication-quality work.

6.2.3 For Research Quality Assurance

AI-simulated expert review achieves goals of human peer review with temporal efficiency:

- **Pre-submission quality control:** Panel review identifies gaps, demands rigor, and validates contributions before venue submission, reducing post-submission revision burden.
- **Systematic improvement:** P1/P2/P3 priority classification focuses revision effort on high-impact issues, with P1 items accounting for 72% of quality improvement [13].
- **Multi-scale synthesis:** Hierarchical review architecture (paper/panel/board tiers) surfaces cross-cutting patterns invisible at individual paper level [12].

6.3 Limitations and Future Work

6.3.1 Study Limitations

This case study has five primary limitations that prevent causal inference:

N=1 treatment group prevents statistical inference: One panel-driven paper vs. three traditional papers is insufficient for hypothesis testing, confidence intervals, or establishing causality. All quantitative comparisons are descriptive observations that generate hypotheses rather than test them.

Selection bias: The VRA compliance paper (panel-driven) addresses a more complex problem than traditional papers (general redistricting). Problem complexity may explain observed differences independent of review methodology. We cannot rule out this confound.

Single author: One researcher (the author, who created the panel system) produced all papers. Multi-user studies with varying expertise levels would test whether patterns transfer to other users.

Domain specificity: Computational redistricting may not represent all domains. Replication across NLP, systems, HCI, and other areas is needed to assess generalizability.

Pre-submission assessment: Quality metrics are assessed before venue submission. Tracking publication outcomes (acceptance rates, citation impact, reproducibility by others) would provide external validation.

6.3.2 Proposed Experimental Design for Future Validation

To rigorously test hypotheses generated by this case study, we propose the following experimental design:

Sample size: Minimum $N=20$ papers per condition (traditional vs. panel-driven) to achieve 80% power for detecting medium effect sizes (Cohen’s $d \geq 0.5$) at $\alpha = 0.05$.

Randomization: For M research questions of matched complexity (assessed by expert raters blind to condition), randomly assign each to traditional or panel-driven workflows. Counterbalance temporal order to control for learning effects.

Multi-user design: Recruit K users ($K \geq 10$) with varying expertise levels (novice, domain expert, methodology expert, both). Assign each user to complete papers in both conditions (within-subjects design) to control for individual differences.

Blind assessment: Have external raters (not involved in paper production) assess quality across 10 dimensions without knowing which condition each paper came from. Compute inter-rater reliability (Krippendorff’s α) and use consensus ratings for analysis.

Ablation studies: Systematically vary:

- Review source (user vs. AI panel vs. human experts)
- Iteration count (1, 2, 4, 8 rounds)
- Panel size (3, 5, 7 reviewers)
- Reviewer selection (domain match vs. random)

Cost tracking: Measure time investment for users, AI compute costs, and iteration overhead to assess efficiency tradeoffs.

Longitudinal outcomes: Submit papers to venues and track acceptance rates, reviewer scores, citation impact at 1 and 3 years post-publication, and reproducibility by independent teams.

Analysis plan: Pre-register hypotheses, analysis methods, and interpretation criteria. Use mixed-effects models to account for nesting (papers within users, users within expertise levels). Report effect sizes with confidence intervals, not just p-values.

This design would provide definitive evidence about whether, when, and for whom panel-driven review improves research quality, moving beyond the hypothesis-generating case study presented here.

6.4 Closing Remarks

The central question motivating this work was: **Does structured peer review simulation improve AI-assisted research quality, or merely add process overhead?**

Our case study observations suggest the former may be true, though definitive evidence requires experimental validation. Panel-driven research demonstrated +127% quality differences by transforming the AI’s role from executor to autonomous researcher responding to iterative expert

review. The panel acts as a scientific driver, systematically questioning assumptions, demanding rigorous testing, embracing negative results, and elevating standards to journal-level rigor.

Critically, breakthrough innovations emerged *because* of panel-driven systematic failure exploration, not despite it. The edge-weighting paradigm shift and 42% threshold theory arose only after 8+ rounds of panel-demanded alternatives exposed fundamental multi-constraint limitations. Traditional research would have documented these limitations and moved on; panel-driven research treated them as puzzles demanding resolution.

This suggests a provocative conclusion: **For AI-assisted research, structured peer review may be more than quality assurance—it may be the primary driver of scientific discovery.** By forcing systematic exploration, embracing failures, and demanding theory building, panel review elevates AI-assisted work from “working methodology” to “publishable science.”

As AI research agents become more capable, the question is not whether they can execute user directions (clearly yes), but whether they can conduct autonomous investigations meeting expert standards. Our work suggests the answer is also yes, provided they operate within structured review frameworks that demand rigor, systematic exploration, and theory building.

The future of AI-assisted research may not be better assistants executing user directions, but autonomous researchers responding to expert review—with users contributing domain expertise and facilitation rather than research direction. This collaboration model, tested here on computational redistricting, may generalize to any domain where breakthrough contributions require systematic failure exploration and expert-level scrutiny.

References

- [1] D. W. Aksnes et al. Citations, citation indicators, and research quality. *Quantitative Science Studies*, 2019.
- [2] Anthropic. Claude sonnet 4.5 technical report, 2025.
- [3] D. Beach and R. B. Pedersen. *Process-Tracing Methods*. University of Michigan Press, 2019.
- [4] D. Boiko et al. Autonomous chemical research with large language models. *Nature*, 2023.
- [5] L. Bornmann. Scientific peer review. *Annual Review of Information Science and Technology*, 2011.
- [6] A. Bran et al. Chemcrow: Augmenting llms with chemistry tools. *Nature Machine Intelligence*, 2024.
- [7] J. Chen et al. Neural approaches to paper review generation. In *Proc. ACL*, 2019.
- [8] Consensus. Consensus: Ai-powered academic search, 2023.
- [9] G. Della-Libera. Ai-simulated expert review: A methodology for pre-submission paper assessment. In *Proc. CHI*, 2026.
- [10] G. Della-Libera. Calibrating ai reviewer personas. In *Proc. EMNLP*, 2026.
- [11] G. Della-Libera. From reviews to revisions: Automated synthesis. In *Proc. AAAI*, 2026.
- [12] G. Della-Libera. Hierarchical review architecture. In *Proc. ICSE*, 2026.
- [13] G. Della-Libera. Multi-round revision dynamics. In *Proc. NeurIPS D&B*, 2026.

- [14] F. Dell’Acqua et al. Navigating the jagged technological frontier. *Harvard Business School Working Paper*, 2023.
- [15] N. Dycke et al. Automated review quality assessment. In *Proc. EMNLP*, 2023.
- [16] Elicit. Elicit: The ai research assistant, 2023.
- [17] S. Firestein. *Failure: Why Science is So Successful*. Oxford University Press, 2015.
- [18] Y. Gao et al. Apo: Automated peer review analysis. *Transactions on Machine Learning Research*, 2023.
- [19] K. Gero et al. Social dynamics of ai idea generation. In *Proc. CSCW*, 2024.
- [20] D. Glass and N. Hall. Hypothesis clarity in philosophical research. *Erkenntnis*, 2021.
- [21] C. J. Lee et al. Bias in peer review. *Journal of the American Society for Information Science and Technology*, 2013.
- [22] M. Lee et al. Coauthor: Designing a human-ai collaborative writing dataset. In *Proc. CHI*, 2022.
- [23] W. Liang et al. Can large language models help researchers search literature? *arXiv preprint*, 2024.
- [24] K. Lu et al. Co-scientist: Ai-assisted hypothesis generation. *Science*, 2024.
- [25] R. D. Riley et al. *Cochrane Handbook for Systematic Reviews*. Wiley, 2019.
- [26] B. Shneiderman. Human-centered ai. *Oxford University Press*, 2022.
- [27] F. Squazzoni et al. Peer review and editorial decision-making at covid-19 speed. *Research Policy*, 2020.
- [28] V. Stodden et al. Enhancing reproducibility for computational methods. *Science*, 2016.
- [29] A. Tomkins et al. Reviewer bias in single- versus double-blind peer review. In *Proc. National Academy of Sciences*, 2017.
- [30] L. Waltman. A review of the literature on citation impact indicators. *Journal of Informetrics*, 2016.
- [31] B. Wang et al. Chemreasoner: Autonomous chemical research. *arXiv preprint*, 2024.
- [32] Q. Wang et al. Scientific discovery in the age of ai. *Nature*, 2023.
- [33] X. Wang et al. Reviewing research papers with neural networks. In *Proc. EMNLP*, 2020.
- [34] X. Wang et al. Can llms replace human collaborators in research? *Nature Human Behaviour*, 2024.
- [35] Y. Zhang et al. Co-author: Human-ai collaborative writing. In *Proc. CHI*, 2023.
- [36] Y. Zhou et al. Large language models in research. *Nature*, 2024.